

# Yolk Manual

Revised November 26th. 2012

# Contents

<b>I</b>	<b>General Information</b>	<b>5</b>
<b>1</b>	<b>Copyright and License</b>	<b>5</b>
<b>2</b>	<b>What is Yolk?</b>	<b>5</b>
2.1	The Yolk demo application . . . . .	5
2.2	The source code . . . . .	5
2.3	Building and installing Yolk . . . . .	6
2.4	The files Yolk depend upon . . . . .	6
2.5	The Yolk packages naming . . . . .	6
<b>II</b>	<b>The Yolk Packages</b>	<b>7</b>
<b>3</b>	<b>Yolk</b>	<b>7</b>
<b>4</b>	<b>Yolk.Cache.Discrete_Keys</b>	<b>7</b>
4.1	The generic formal parameters . . . . .	7
4.2	Instantiation . . . . .	7
4.3	Writing to the cache . . . . .	8
4.4	Reading from the cache . . . . .	8
4.5	Checking if a key is valid . . . . .	9
4.6	Clearing keys and the entire cache . . . . .	9
4.7	Cleanup - Getting rid of stale elements . . . . .	9
<b>5</b>	<b>Yolk.Cache.String_Keys</b>	<b>9</b>
5.1	The generic formal parameters . . . . .	10
5.2	Instantiation . . . . .	10
5.3	Writing to the cache . . . . .	10
5.4	Reading from the cache . . . . .	11
5.5	Checking if a key is valid . . . . .	11
5.6	Clearing keys and the entire cache . . . . .	11
5.7	How much is in there? . . . . .	12
5.8	Cleanup - Keeping cache size in check . . . . .	12

<b>6</b>	<b>Yolk.Config_File_Parser</b>	<b>12</b>
6.1	The generic formal parameters . . . . .	13
6.2	Exceptions . . . . .	13
6.3	Instantiation . . . . .	13
6.4	Re-loading configuration files . . . . .	14
6.5	Getting values . . . . .	14
6.6	Checking if a KEY has a VALUE . . . . .	15
<b>7</b>	<b>Yolk.Configuration</b>	<b>15</b>
7.1	Get the AWS specific configuration settings . . . . .	16
<b>8</b>	<b>Yolk.Email</b>	<b>16</b>
8.1	Exceptions . . . . .	16
8.2	The Yolk.Email types . . . . .	16
<b>9</b>	<b>Yolk.Email.Composer</b>	<b>18</b>
9.1	Building and sending an email, the easy way . . . . .	18
9.2	Building and sending email, the hard way . . . . .	20
<b>10</b>	<b>Yolk.Handlers</b>	<b>21</b>
<b>11</b>	<b>Yolk.Log</b>	<b>21</b>
<b>12</b>	<b>Yolk.Not_Found</b>	<b>22</b>
<b>13</b>	<b>Yolk.Process_Control</b>	<b>22</b>
13.1	Exceptions . . . . .	23
13.2	Using Yolk.Process_Control . . . . .	23
<b>14</b>	<b>Yolk.Process_Owner</b>	<b>24</b>
14.1	Exceptions . . . . .	24
14.2	Using Yolk.Process_Owner . . . . .	24
<b>15</b>	<b>Yolk.Server</b>	<b>25</b>
15.1	Yolk.Server.Create . . . . .	25
15.2	Yolk.Server.Start . . . . .	25
15.3	Yolk.Server.Stop . . . . .	26

<b>16 Yolk.Static_Content</b>	<b>26</b>
16.1 Yolk.Static_Content.Static_Content_Cache_Setup . . . . .	27
16.2 Yolk.Static_Content.Compressable . . . . .	27
16.3 Yolk.Static_Content.Non_Compressable . . . . .	29
16.4 Yolk.Static.Set_Cache_Options . . . . .	30
<b>17 Yolk.Syndication</b>	<b>30</b>
17.1 Exceptions . . . . .	31
17.2 The Yolk.Syndication types . . . . .	31
17.3 Yolk.Syndication.New_Atom_Feed . . . . .	32
17.4 Yolk.Syndication.New_Atom_Entry . . . . .	33
17.5 Yolk.Syndication.New_Atom_Entry_Source . . . . .	33
17.6 Yolk.Syndication.Writer . . . . .	33
17.7 Counting the amount of entries in a feed . . . . .	34
17.8 Clearing and deleting entries . . . . .	34
17.9 Getting the Atom feed . . . . .	35
<b>18 Yolk.Utilities</b>	<b>36</b>
<b>19 Yolk.Whoops</b>	<b>36</b>

## Part I

# General Information

## 1 Copyright and License

This document is copyright (C) 2010-2012, Thomas Løcke. You may copy this document, in whole or in part, in any form or by any means, as is or with alterations, provided that (1) alterations are clearly marked as alterations and (2) this copyright notice is included unmodified in any copy.

Yolk is GPLv3 software. You should have received a copy of the GNU General Public License and a copy of the GCC Runtime Library Exception along with this program; see the files COPYING3 and COPYING.RUNTIME respectively. If not, see <http://www.gnu.org/licenses/>.

## 2 What is Yolk?

Yolk is a collection of packages that aim to help build solid web-applications using Ada. Yolk itself doesn't do a whole lot that can't be accomplished simply by using AWS and the GNAT Component Collection (GNATcoll), but it does make the job of building complete web-applications a bit simpler. Things like changing user for the running application, catching POSIX signals such as SIGKILL, sending log data to syslogd, adding basic static content handlers, creating and starting/stopping an AWS powered HTTP server and building Atom syndication XML are all made a bit easier with Yolk.

A Yolk application is in reality an AWS application, with some sugar added, so you're not really building a Yolk web-application, as much as you're building an AWS web-application. What I'm getting at, is that you need to understand how to use AWS, in order for Yolk to make any kind of sense. What you get when using Yolk is the little things that AWS does not readily provide.

### 2.1 The Yolk demo application

Reading this manual will of course (I hope!) help you understand how to use Yolk, but please consider taking a closer look at the Yolk demo application to get a feel for how Yolk is actually used. The demo is heavily commented, so it should be fairly easy to understand what's going on. The demo application is also very suitable as a foundation for other AWS/Yolk applications.

It is much easier to show how to use Yolk, than it is to write down all possible usage scenarios. With the combination of this manual, the Yolk source files and the demo application, you should be able to make full use of the Yolk packages in your own applications.

### 2.2 The source code

The Yolk source code is the best documentation there is. This document is never going to be as comprehensive as the actual source, so I'll strongly suggest having the source

code available as you read this document. What you will find in this document are short descriptions of what a package is meant to do and perhaps small usage examples, not a complete rundown of every type and procedure in a package.

## 2.3 Building and installing Yolk

See the README and INSTALL files. These are found in the Yolk root directory.

## 2.4 The files Yolk depend upon

When you read this document and the Yolk source code, you'll notice that quite a few packages depend on various files being available at specified locations. This is for example the case with the *Yolk.Whoops* package that expects its template file to be found at the path *templates/system/500.tmpl*

All such "dependencies" will of course be noted accordingly as we go along, but instead of forgetting one or more in your own application, I'd much rather encourage using the demo application as a foundation for your own applications, since all these fixed paths and files has been properly added to the demo.

I also recommend compiling and running the demo, to make sure your Yolk install is working as intended. Just read the *demo/README* and *demo/INSTALL* files for instructions on how to get it up and running.

## 2.5 The Yolk packages naming

The Yolk packages are pretty diverse, ranging from process control to sending email. I've tried naming them as sensibly as possible, in the hope that the package names alone give away their function. If I've failed, well, you're just going to have to refer to this document or take a look at the source for yourself.

## Part II

# The Yolk Packages

### 3 Yolk

The Yolk main package currently only contain only a few things: The Yolk *Version* string and a *Config\_File* function to get the location of the configuration file. These are used in a few places, for example in the *directory.tmpl* template file (the version string) and in the *Yolk.Configuration* package (the *Config\_File* function).

### 4 Yolk.Cache.Discrete\_Keys

If a piece of data doesn't change very often and it is expensive to build, then caching it might be worthwhile. Instead of going to a file or database on every hit, you simply go to the cache and grab the latest version from there. This is **very** fast, at the cost of some memory.

If you know exactly what you want to cache, the *Yolk.Cache.Discrete\_Keys* package might be just what you need.

#### 4.1 The generic formal parameters

These are:

```
generic
  type Key_Type is (<>);
  type Element_Type is private;
  Max_Element_Age : Duration := 3600.0;
package Yolk.Cache.Discrete_Keys is
  ...
```

The *Max\_Element\_Age* defaults to one hour. You should obviously set this to whatever suits your needs. This timer is used for all content in the cache. You cannot set this individually for each element.

#### 4.2 Instantiation

If for example we have two different sets of data (Foo and Bar) that are expensive to build, we can instantiate a *Discrete\_Keys* package to handle this:

```
type Cache_Keys is (Foo, Bar);
package My_Cache is new Yolk.Cache.Discrete_Keys
  (Key_Type      => Cache_Keys,
   Element_Type => Unbounded_String);
```

And that is all. We now have a *My\_Cache* object that can hold two objects: *Foo* and *Bar*. These are of the type *Unbounded\_String* and they have a *Max\_Element\_Age* of 3600.0 seconds.

### 4.3 Writing to the cache

Before we can read something from the cache, we must first write something to it:

```
declare
  Foo_Value : Unbounded_String := To_Unbounded_String ("Foo");
begin
  My_Cache.Write (Key    => Foo,
                  Value => Foo_Value);
end;
```

That is all it takes: “Foo” is now safely tucked away in the *My\_Cache* object, and will be so for 3600.0 seconds. Calling *Write* with the *Foo* key will always overwrite earlier written *Foo* elements, no matter their age.

### 4.4 Reading from the cache

A cache obviously only makes sense if you intend to read from it. In our case we want to get our hands on the previously written “Foo” value:

```
declare
  Valid : Boolean := False;
  Value : Unbounded_String;
begin
  My_Cache.Read (Key      => Foo,
                 Is_Valid => Valid,
                 Value    => Value);
  if Valid then
    — do something interesting with the data
  else
    — the Foo data is invalid.
  end if;
end;
```

In order for an element to be valid (the *Is\_Valid* parameter is true), it must:

1. have been added to the cache in the first place
2. be younger than *Max\_Element\_Age*

If *Is\_Valid* is *False*, then *Value* is undefined. Note that if *Is\_Valid* is *False* then *Key* is removed from the cache, if it exists.



## 4.5 Checking if a key is valid

If you need to check whether a specific key exists in the cache and is valid, then you must use the *Is\_Valid* function.

```
if My_Cache.Is_Valid (Foo) then
  — Foo is good!
else
  — Foo is bad!
end if ;
```

This follows the same rules as the *Is\_Valid* parameter for the *Read* procedure.

## 4.6 Clearing keys and the entire cache

For clearing of keys and the entire cache we have, naturally, two *Clear* procedures:

```
— First we clear the Foo key
My_Cache.Clear (Key => Foo);

— And then we clear the entire cache
My_Cache.Clear ;
```

And that's all it takes.

## 4.7 Cleanup - Getting rid of stale elements

Calling Cleanup will delete all stale elements from the cache:

```
My_Cache.Cleanup ;
```

Note that this is potentially a very expensive operation if the cache is large, as the entire cache is iterated and every element tested for its age. Use with care.

# 5 Yolk.Cache.String\_Keys

This package is almost similar to the *Yolk.Cache.Discrete\_Keys* package. The biggest difference is that where the *Discrete\_Keys* cache package requires that you define a type for the keys, this package use regular *String* as keys.

The implications of this difference between the two cache packages are subtle. Both have the same *Read*, *Write*, *Is\_Valid* and *Clear* procedures and functions, so in that sense the two packages are the same. The biggest difference lies in the available generic formal parameters and the functionality of the *Cleanup* procedure.

## 5.1 The generic formal parameters

These are:

```
generic
  type Element_Type is private;
  Cleanup_Size      : Positive := 200;
  Cleanup_On_Write  : Boolean   := True;
  Max_Element_Age   : Duration  := 3600.0;
  Reserved_Capacity : Positive := 100;
package Yolk.Cache.Discrete_Keys is
  ...
```

When the amount of elements in the cache  $\geq$  *Cleanup\_Size*, then the *Cleanup* procedure is called by *Write*, if *Cleanup\_On\_Write* is set to Boolean *True*. *Cleanup\_Size* is a sort of failsafe for this cache package. Since we can't know for sure what is being added (we don't know the keys beforehand), we need to make sure it doesn't gobble up all available resources. Set this number high enough that it'll never trigger under normal circumstances, but low enough that it'll prevent resource exhaustion in case of errors.

The *Max\_Element\_Age* defaults to one hour. You should obviously set this to whatever suits your needs. This timer is used for all content in the cache. You cannot set this individually for each element.

*Reserved\_Capacity* should be set as close as possible to the expected final size of the cache. If your best guesstimate is 200 elements in the cache, then set this to 200. Note that this setting has no bearing on the actual size of the cache. The cache will happily grow beyond the *Reserved\_Capacity* value.

## 5.2 Instantiation

Instantiating *String\_Keys* is done like this:

```
package My_Cache is new Yolk.Cache.String_Keys
  (Element_Type      => Unbounded_String ,
   Reserved_Capacity => 200);
```

And that is all. We now have a *My\_Cache* object that can hold objects of the type *Unbounded\_String*, all of which have a *Max\_Element\_Age* of 3600.0 seconds. Also we've told the cache to set aside at least 200 positions for content.

## 5.3 Writing to the cache

Before we can read something from the cache, we must first write something to it:

```
declare
  Value : Unbounded_String := To_Unbounded_String ("42");
begin
  My_Cache.Write (Key   => "Foo" ,
                  Value => Value);
end;
```

“42” is now safely tucked away in the *My\_Cache* object under the key “Foo”, and will be so for 3600.0 seconds. Calling *Write* with the “Foo” String will always overwrite earlier written “Foo” elements, no matter their age.

## 5.4 Reading from the cache

A cache obviously only makes sense if you intend to read from it. In our case we want to get our hands on the previously written “Foo” value:

```
declare
  Valid : Boolean := False;
  Value : Unbounded_String;
begin
  My_Cache.Read (Key      => "Foo",
                  Is_Valid => Valid,
                  Value    => Value);
  if Valid then
    — do something interesting with the data
  else
    — the Foo data is invalid.
  end if;
end;
```

In order for an element to be valid (the *Is\_Valid* parameter is true), it must:

1. have been added to the cache in the first place
2. be younger than *Max\_Element\_Age*

If *Is\_Valid* is *False*, then *Value* contains undefined garbage. Note that if *Is\_Valid* is *False* then *Key* is removed from the cache, if it exists.

## 5.5 Checking if a key is valid

If you need to check whether a specific key exists in the cache and is valid, then you need to use the *Is\_Valid* function.

```
if My_Cache.Is_Valid ("Foo") then
  — Foo is good!
else
  — Foo is bad!
end if;
```

This follows the same rules as the *Is\_Valid* parameter for the *Read* procedure.

## 5.6 Clearing keys and the entire cache

For clearing of keys and the entire cache we have, naturally, two *Clear* procedures:

```
-- First we clear the Foo key
My_Cache.Clear (Key => "Foo");

-- And then we clear the entire cache
My_Cache.Clear;
```

## 5.7 How much is in there?

With the *Discrete\_Keys* cache we obviously always know the exact amount of keys available, since we've defined the keys ourselves. This is not the case with the *String\_Keys* cache, where any *String* can be a key. If we need to know how many elements that are currently in the cache, we call the *Length* function:

```
if My_Cache.Length > 1000 then
  -- Woa! Lots of stuff in the cache..
end if;
```

Note that *Length* count both valid and invalid elements.

## 5.8 Cleanup - Keeping cache size in check

if *Cleanup\_On\_Write* is *True*, then *Cleanup* is called by *Write* whenever the size of the cache reach *Cleanup\_Size*. It is of course also possible to call it manually:

```
if My_Cache.Length > 1000 then
  My_Cache.Cleanup;
end if;
```

If you've set *Cleanup\_On\_Write* to Boolean *False* and the *String* keys are coming from outside sources, then you really should make sure you call *Cleanup* on a regular basis.

## 6 Yolk.Config\_File\_Parser

This package enables you to access KEY/VALUE pairs in configuration files that are written in the style:

```
# This is a comment
-- This is also a comment
KEY VALUE
```

Keys are case-insensitive, so *FOO*, *foo* and *fOo* are all the same. Blank lines and comments are ignored and so is pre/postfixed whitespace. It is not necessary to quote values that contain whitespace, to this:

```
KEY some value with whitespace
```

is perfectly valid, and will return “*some value with whitespace*” when calling *Get (KEY)*. If *VALUE* is *Boolean True* or *False* (case-insensitive), then the *KEY* can be returned as a *String* or a *Boolean*, depending on the target type. If the target type does not match the *VALUE* and no sensible conversion can be made, then a *Conversion\_Error* exception is raised. No dummy values are returned at any time.

To clear a default value, simply add the key to the configuration file, with no value set.

## 6.1 The generic formal parameters

These are:

```
generic
  use Ada.Strings.Unbounded;
  type Key_Type is (<>);
  type Defaults_Array_Type is array (Key_Type) of Unbounded_String;
  Defaults      : in Defaults_Array_Type;
  Config_File  : in String;
package Yolk.Config_File_Parser is
  ...
```

*Config\_File* is of course the name and location of the configuration file.

## 6.2 Exceptions

There are 3 different exceptions that can be raised by the *Yolk.Config\_File\_Parser* package. These are:

- *Unknown\_Key*. This is raised if an unknown key has been found in the configuration file given when instantiating the package or when *Load\_File* is called.
- *Cannot\_Open\_Config\_File*. This is raised when *Config\_File* cannot be read.
- *Conversion\_Error*. This is raised when a value cannot be converted to the target type, ie. the value “42” to a *Boolean*.

## 6.3 Instantiation

*Yolk.Config\_File\_Parser* is a generic package, so in order to use it, you have to instantiate it, like this:

```

package My_Configuration is

    type Keys is (Foo, Bar);
    type Defaults_Array is array (Keys) of Unbounded_String;

    Default_Values : constant Defaults_Array :=
        (Foo => To_Unbounded_String ("some foo"),
         Bar => To_Unbounded_String ("some bar"));

    package Config is new Yolk.Config_File_Parser
        (Key_Type => Keys,
         Defaults_Array_Type => Defaults_Array,
         Defaults => Default_Value,
         Config_File => "config.ini");

end My_Configuration;

```

Here we instantiate the *Config* package with *config.ini* as the configuration file. This means that KEY/VALUE pairs found in this file will overwrite the default values set in the *Default\_Values* array. Setting a default value to *Null\_Unbounded\_String* means the value is empty.

Note that the *config.ini* file does not have to contain all the valid keys. It is perfectly fine to only add those keys that have non-default values to the configuration file.

## 6.4 Re-loading configuration files

With the *Load\_File* procedure you can re-load a new configuration file into your *Config* package:

```

My_Configuration.Config.Load_File ("new_config.ini");

```

Now the KEY/VALUE pairs of *new\_config.ini* will overwrite the ones originally found in the *config.ini* file the package was instantiated with. You can do this as many times as you like. Note that you cannot change what KEY's are valid, so if the *new\_config.ini* file contains unknown keys, *Load\_File* will raise the *Unknown\_Key* exception.

## 6.5 Getting values

With instantiation and loading of configuration files out of the way, it is now time to get to the configuration values. To get the value of the *Foo* key, you do:

```

My_Configuration.Config.Get (Foo);

```

There are Get functions for the following types:

- *Boolean*
- *Duration*

- *Float*
- *Integer*
- *String*
- *Unbounded\_String*

Empty keys simply return an empty *String* or a *Null\_Unbounded\_String*, depending on the target type. If a key is empty and the target type is not a *String* or an *Unbounded\_String*, then the *Conversion\_Error* exception is raised.

## 6.6 Checking if a KEY has a VALUE

You can check if a key has a value with the *Has\_Value* function:

```
if Has_Value (Foo) then
  Put_Line ("Foo has a value");
end if;
```

Basically all this function does is return *Boolean True* if the value of the given key is not a *Null\_Unbounded\_String*.

## 7 Yolk.Configuration

This package is a bit of an oddball, as all it does is instantiate the *Yolk.Config\_File\_Parser* generic with the default AWS and Yolk configuration values. This is used by Yolk internally, but also by the AWS component of your application. The instantiation looks like this:

```
package Config is new Config_File_Parser
(Key_Type           => Keys,
 Defaults_Array_Type => Defaults_Array,
 Defaults           => Default_Values,
 Config_File        => Config_File);
```

The *Config\_File* function call return either the default Yolk configuration file (*configuration/-config.ini*) or a user specified configuration file given by the *--yolk-config-file* command line argument, so starting for example the Yolk demo like this:

```
./yolk_demo --yolk-config-file /etc/yolk-config.ini
```

will force the demo to look for the */etc/yolk-config.ini* configuration file.

There's a fully commented *config.ini.dist* file available in the *extras/* directory. I recommend taking a look at the Yolk demo application to see how the *Yolk.Configuration* package is used.

## 7.1 Get the AWS specific configuration settings

On some occasions it might be necessary to get the AWS configuration object. This can easily be accomplished by calling the *Get\_AWS\_Configuration* function:

```
AWS_Config : constant AWS.Config.Object :=  
  Yolk.Configuration.Get_AWS_Configuration ;
```

## 8 Yolk.Email

Using *Yolk.Email* and the child package *Yolk.Email.Composer* you can build and send more or less any kind of email:

- Plain text
- Multipart/Alternative
- Multipart/Mixed

The package supports adding multiple SMTP servers, meaning you can add as many as you need, and the email will then be send via the first one that accepts it.

The *Yolk.Email* package define 4 exceptions and 3 types. The facilities for actually constructing and sending the email are found in *Yolk.Email.Composer*.

### 8.1 Exceptions

These are:

- *Attachment\_File\_Not\_Found*. Is raised if a file attachment is not found at the given path.
- *No\_Address\_Set*. Is raised if the address component of a To, Reply-To, From, Bcc/Cc header is missing.
- *No\_Sender\_Set\_With\_Multiple\_From*. Is raised when an email contains multiple From headers but no Sender header, as per RFC-5322, 3.6.2. <http://tools.ietf.org/html/rfc5322>
- *No\_SMTP\_Host\_Set*. Is raised if the SMTP host list is empty, ie. no SMTP host has been set for sending the email.

### 8.2 The Yolk.Email types

When using *Yolk.Email.Composer* to build and send emails, three types declared in *Yolk.Email* are central:

1. *Character\_Set*



## 2. *Recipient\_Kind*

## 3. *Structure*

The *Character\_Set* type define what character set is used when data is added to an email *Structure* object. For example looking at the *Yolk.Email.Composer.Add\_From* procedure, we see that the *Charset* parameter defaults to *US\_ASCII*:

```
procedure Add_From
  (ES      : in out Structure ;
   Address : in      String ;
   Name    : in      String := "";
   Charset : in      Character_Set := US_ASCII);
```

This does not mean that *Yolk.Email.Composer.Add\_From* will encode *Name* as *US\_ASCII*, instead it means that the data given in *Name* already is *US\_ASCII*. So if *Name* had contained an ISO-8859-1 encoded *String*, then the call would've looked like this:

```
declare
  use Yolk.Email ;

  Email : Structure ;
begin
  Composer.Add_From
    (ES      => Email ,
     Address => "alice@domain.tld" ,
     Name    => "Alice" ,
     Charset => ISO_8859_1);
end;
```

In this case you will end up with a From header looking like this:

```
From: =?ISO-8859-1?Q?Alice?= <alice@domain.tld>
```

So bear in mind that it is your responsibility to ensure that your data, and the *Character\_Set* parameter match.

The *Recipient\_Kind* type define the kind of recipient that is being added to an email. If you've worked with email, these three should be familiar to you:

1. *Bcc*
2. *Cc*
3. *To*

When adding recipients to an email *Structure* the default is *To*, but since not all recipients are equal, you can change the kind of recipient to *Bcc* or *Cc*, according to your needs.

The *Structure* type is at the core of it all. You declare an object to be of the *Structure* type, and then you use the *Yolk.Email.Composer* facilities to build and send the email.

## 9 Yolk.Email.Composer

The actual tools for building and sending an email is found in this package. Here are tools for building emails from the ground up and there are a few convenience procedures if you just need to send a simple email with no bells or whistles.

I'm not going to go through ever procedure in this package, instead I'll show an example on how to build an email from the ground up and how to use one of the convenience procedures.

### 9.1 Building and sending an email, the easy way

There are two convenience procedures in *Yolk.Email.Composer* for sending emails without having to do a whole lot of work/thinking. They are both named *Send* and they look like this:

```
procedure Send
(ES
  From_Address : in out String;
  From_Name    : in    String := "";
  To_Address   : in    String;
  To_Name      : in    String := "";
  Subject      : in    String;
  Text_Part    : in    String;
  SMTP_Server  : in    String := "localhost";
  SMTP_Port    : in    Positive := 25;
  Charset      : in    Character_Set := US_ASCII);

procedure Send
(ES
  From_Address : in out String;
  From_Name    : in    String := "";
  To_Address   : in    String;
  To_Name      : in    String := "";
  Subject      : in    String;
  Text_Part    : in    String;
  HTML_Part    : in    String;
  SMTP_Server  : in    String := "localhost";
  SMTP_Port    : in    Positive := 25;
  Charset      : in    Character_Set := US_ASCII);
```

As you can see, the only difference between these two is that the first one sends plain text emails, while the second one sends *multipart/alternative* with both plain text and HTML parts. Usage is as simple as:

```

declare
  use Yolk.Email;

  Email : Structure;
begin
  Composer.Send (ES          => Email ,
                  From_Address => "alice@domain.tld" ,
                  From_Name    => "Alice" ,
                  To_Address   => "bob@domain.tld" ,
                  To_Name      => "Bob" ,
                  Subject       => "Is this thing on?" ,
                  Text_Part     => "Hey you!" ,
                  Charset       => ISO_8859_1);

  if Composer.Is_Send (Email) then
    — Success!
  else
    — Failure!
  end if;
end;

```

It is possible, and allowed, to call some of the various other procedures prior to calling one of the convenience procedures. If for example you want to add a custom header, it can be done like this:

```

declare
  use Yolk.Email;

  Email : Structure;
begin
  Composer.Add_Custom_Header (ES          => Email ,
                              Name        => "User-Agent" ,
                              Value       => "My User Agent");

  Composer.Send (ES          => Email ,
                  From_Address => "alice@domain.tld" ,
                  From_Name    => "Alice" ,
                  To_Address   => "bob@domain.tld" ,
                  To_Name      => "Bob" ,
                  Subject       => "Is this thing on?" ,
                  Text_Part     => "Hey you!" ,
                  Charset       => ISO_8859_1);

  if Composer.Is_Send (Email) then
    — Success!
  else
    — Failure!
  end if;
end;

```

And with that, the header *User-Agent*: is now added to the email:

User-Agent: My User Agent

It hardly gets any easier than that. Lets move on and see how the above is accomplished the hard way.

## 9.2 Building and sending email, the hard way

It is possible to build an email from the ground up, which obviously allows for a more fine grained control over what is added. It is also a bit more complicated, but not much. Lets try and mimick the easy examples, the “hard” way:

```
declare
  use Yolk.Email;

  Email : Structure;
begin
  Composer.Add_Custom_Header (ES    => Email,
                              Name  => "User-Agent",
                              Value => "My User Agent");

  Composer.Add_From (ES    => Email,
                    Address => "alice@domain.tld",
                    Name   => "Alice",
                    Charset => ISO_8859_1);

  Composer.Add_Recipient (ES    => Email,
                         Address => "bob@domain.tld",
                         Name    => "Bob");

  Composer.Set_Subject (ES    => Email,
                      Subject => "Is this thing on?");

  Composer.Set_Text_Part (ES    => Email,
                        Part => "Hey you!");

  Composer.Add_SMTP_Server (ES    => Email,
                          Host => "localhost");

  Composer.Send (ES => Email);

  if Composer.Is_Send (Email) then
    — Success!
  else
    — Failure!
  end if;
end;
```

Harder yes, but really not all that much more difficult.

## 10 Yolk.Handlers

Most web applications will need to handle static content, such as PNG, HTML and CSS files. *Yolk.Handlers* helps you accomplish that, so you don't have to build your own handlers for these kinds of files.

The following filetypes are supported by *Yolk.Handlers*:

- CSS
- GIF
- HTML
- ICO
- JPG
- JS
- PNG
- SVG
- XML
- XSL

The filetypes that are textual, are compressed according to the *Yolk.Configuration* parameter *Compress\_Static\_Content* which defaults to *False*. The regular expressions for identifying these filetypes are also defined in *Yolk.Configuration* by the *Handler\_\** parameters. These regular expressions are registered by the *AWS.Services.Dispatchers.URI.Register\_Regexp* procedure.

There's only one procedure in the *Yolk.Handlers* package:

```
procedure Set (RH : out AWS.Services.Dispatchers.URI.Handler);
```

You can see an example on how this is used in the demo file *my\_handlers.adb*. There's really very little reason not to use this package for handling of static content, but it is of course not mandatory.

This package makes use of the *Yolk.Static\_Content* package for the actual delivery of the content to the user.

## 11 Yolk.Log

This package serves two purposes:

1. It contains the two callback procedures used to write AWS logging data (access and error) to syslogd.

2. It creates the SQL, SQL\_Cache, SQL\_Error, SQL\_Select, Alert, Critical, Debug, Emergency, Error, Info, Notice and Warning trace handles, and activates them according to the values defined in the configuration file.

Out of the box, a Yolk application requires a running syslogd daemon, as all log data is sent to syslogd. If for some reason you don't want to use syslogd, you're going to have to hack the *Yolk.Log* package, or remove it entirely.

The two procedures named *AWS\_\** are used by the AWS HTTP(S) server. These should not be used for anything else - or rather: If you use them for anything else, the *Message* given is going to be written to syslogd with either the AWS access log label or AWS error log label. There's absolutely no harm in this, except it might be a bit confusing when reading the log data.

The *Trace* procedure is the one that you will be using in your application. You can send log data to the trace handles defined in *Yolk.Log.Trace\_Handles*. All log data is then sent to the syslogd daemon using the facilities set in the configuration file for the given trace handle.

Using *Trace* is about as easy as it gets:

```
Yolk.Log.Trace (Handle => Error ,  
               Message => "Secret sauce to Error!");
```

If you haven't activated a trace handle, then calling *Trace* for that handle does nothing, ie. you don't have to remove all your *Trace* calls from the code if you don't use them.

Note that on many systems the Emergency messages are written to console, so be cautious about using this specific level, unless you've made certain that you can get to the messages.

## 12 Yolk.Not\_Found

The job of this package is to return a HTTP 404 status code and an accompanying simple *not found* HTML page. Its sole function *Generate* is about as simple as they come:

```
function Generate  
  (Request : in AWS.Status.Data)  
  return AWS.Response.Data ;
```

It relies on the template file *demo/exe/templates/system/404.tmpl* to generate the generic 404 HTML page, so if you want to use *Yolk.Not\_Found* in your own application, then remember to bring along this file. Where the *404.tmpl* is placed is defined in the configuration parameter *System\_Templates\_Path*.

Also worth noting is that the *Yolk.Not\_Found.Generate* function is used as the default callback in the demo application. This means that all requested resources that doesn't match a registered dispatcher, is served by *Yolk.Not\_Found.Generate* ie. a 404 is returned. See the *demo/src/my\_handlers.adb* file for more information.

## 13 Yolk.Process\_Control

With *Yolk.Process\_Control* you get the ability to control your application using the SIGINT, SIGPWR and SIGTERM signals. You can also get a PID file placed next to your executable

(see the *PID\_File* configuration parameter), which coupled with the *extras/rc.yolk* script gives you a rather nice and simple way of starting and stopping your application.

## 13.1 Exceptions

These are:

- *Cannot\_Create\_PID\_File*. Is raised if the *PID\_File* cannot be created, eg. if the application lacks permissions to write to the directory where the *PID\_File* is located.
- *Cannot\_Delete\_PID\_File*. Is raised if the *PID\_file* cannot be deleted, eg. if the application lacks permissions to write to the directory where the *PID\_File* is located, or to the *PID\_File* itself.
- *PID\_File\_Exists*. Is raised when the *PID\_File* already exists, ie. the application is already running or it was shutdown incorrectly.

## 13.2 Using Yolk.Process\_Control

When you use the *Yolk.Process\_Control* package the *Unreserve\_All\_Interrupts* pragma is used. This means that depending on the compiler used one or more interrupt signals may be affected. In the case of the GNAT compiler, this is specifically mentioned in the source of the *Ada.Interrupts.Names* package:

- The pragma *Unreserve\_All\_Interrupts* affects the following signal(s):
- SIGINT: made available for Ada handler

Since neither SIGPWR or SIGTERM are reserved by the compiler, the *Yolk.Process\_Control* package is able to assume control of these signals. You can read more about the pragma *Unreserve\_All\_Interrupts* here. If you compile Yolk with a different compiler than GNAT, then please check if one of the affected signals are reserved.

There are two procedures in the *Yolk.Process\_Control* package:

```
procedure Stop ;  
  
procedure Wait ;
```

When you call the *Wait* procedure, you basically hang there until

1. The *Stop* procedure is called
2. The application receives a SIGINT, SIGPWR or SIGTERM signal

This is quite handy for applications, that need some sort of loop to keep them from terminating. You can see an example on how this can be done in the *demo/src/yolk\_demo.adb* file.

When *Wait* is called, subsequent calls to *Wait* are ignored, unless a call to *Stop* has been made or the application has received one of the SIGINT, SIGPWR or SIGTERM signals. So it's perfectly valid to do:

```
Wait;  
— Stop called from somewhere in the app  
— Do something...  
Wait;  
— The app receives a SIGINT signal  
— Do something...  
Wait;
```

Whether or not this is actually useful I don't know, but it is possible.

## 14 Yolk.Process\_Owner

When it is necessary to change the owner of a process, the *Yolk.Process\_Owner* package is the solution. Obviously this can also be done when starting the application, using various shell tricks, but I find it much cleaner to just let the application handle it by itself.

### 14.1 Exceptions

There's only one:

1. *Username\_Does\_Not\_Exist*. This is raised if the given username doesn't exist on the system.

### 14.2 Using Yolk.Process\_Owner

There's only a single procedure in this package and its specification looks like this:

```
procedure Set_User  
  (Username : in String);  
  — Set the process owner to Username.
```

Please note that when changing the user ID of the application with *Set\_User*, the group ID is changed to the first group the given user is a member of.

Usage is as simple as expected:

```
declare  
begin  
  Set_User (Username => "billybob");  
exception  
  when Username_Does_Not_Exist =>  
    — User is missing. Do something!  
end;
```

In the file *demo/src/yolk\_demo.adb* you'll find that *Yolk.Process\_Owner.Set\_User* is used in conjunction with the *Yolk.Configuration.Yolk\_User* parameter.



## 15 Yolk.Server

This is a convenience package to handle creating, starting and stopping an AWS HTTP server. The biggest drawback to this package is that it can only create and manage *one* AWS server.

Note a potential manual operation when using the *Yolk.Server* package : Setting the *Cache-Control* header. This can become necessary if you

1. use *Yolk.Server* to create and start your AWS server AND
2. use the *Yolk.Handlers* package to register dispatchers for static content AND
3. have the *Compress\_Static\_Content* configuration parameter set to *True* AND
4. aren't content with the default *Cache-Control* header

In that case you can use the *Yolk.Static\_Content.Set\_Cache\_Options* procedure to work your own magic on the *Cache-Control* header sent to clients requesting static content. See the *Yolk.Static\_Content* chapter for more information.

### 15.1 Yolk.Server.Create

Creating a server is done by calling the *Create* function. This function accepts one parameter that should be fairly self-explanatory.

```
type HTTP is tagged limited private ;

function Create
  (Unexpected : in AWS.Exceptions.Unexpected_Exception_Handler)
  return HTTP;
```

### 15.2 Yolk.Server.Start

```
procedure Start
  (WS           : in out HTTP;
   Dispatchers  : in      AWS.Dispatchers.Handler'Class);
```

This does as expected: Starts the server with *Dispatchers*.

When calling *Start* several things happen:

1. If the configuration parameter *Load\_MIME\_Types\_File* is *True* then the MIME types file given in *MIME\_Types\_File* is loaded into AWS.
2. If the configuration parameter *Start\_WebSocket\_Servers* is *True* then the AWS Web-Socket servers are started.
3. *Yolk.Static\_Content.Static\_Content\_Cache\_Setup* is called with its default parameters if the configuration parameter *Compress\_Static\_Content* is *True*.

4. If session support is turned on for the server and the *Session\_Data\_File* exists, then the session data is loaded from this file.
5. The *Yolk.Log.AWS\_Access\_Log\_Writer* procedure is registered to handle all AWS access log data if the configuration parameter *AWS\_Access\_Log\_Activate* is *True* .
6. The *Yolk.Log.AWS\_Error\_Log\_Writer* procedure is registered to handle all AWS error log data if the configuration parameter *AWS\_Error\_Log\_Activate* is *True* .

### 15.3 Yolk.Server.Stop

```
procedure Stop  
(WS : in out HTTP);
```

This does as expected: Stops everything that was started when the *Yolk.Server.Start* call was made. Note that if sessions support is turned on for the server, then the *Dispatchers* are switched for a plain *Not Found* dispatcher when *Stop* is called. This is done to prevent session tampering while the server is shutting down.

## 16 Yolk.Static\_Content

Most web applications have a lot of static content, ie. stuff that rarely changes. Things like PNG's, HTML, CSS and Javascript. These are content types that are common for most websites, so a application is going to have to handle these in some way. This is where *Yolk.Static\_Content* comes in. Two kinds of files are handled by *Yolk.Static\_Content*:

- Compressable (XML, HTML, CSS, JS and so on)
- Non-compressable (PNG, JPG, GIF, ICO and so on)

It is up to you, the user, to decide whether a specific kind of file is compressable or not - the package does not make any such assumptions. The difference between the two, is that compressable files are compressed prior to delivery, if the clients HTTP request includes a *Accept-Encoding: gzip* header. Non-compressable files are simply returned as is. For both kinds, a generic 404 is returned if the requested file doesn't exist.

*Yolk.Static\_Content* is affected by five configuration settings:

- *Compress\_Static\_Content*
- *Compress\_Static\_Content\_Minimum\_File\_Size*
- *Compressed\_Static\_Content\_Cache*
- *Compressed\_Static\_Content\_Max\_Age*
- *WWW\_Root*

You should carefully read the *demo/exe/configuration/config.ini* file for information on what exactly these do.

## 16.1 `Yolk.Static_Content.Static_Content_Cache_Setup`

The configuration parameter *Compressed\_Static\_Content\_Cache* defines where the compressed version of the static content is saved. When your application is started, this directory may or may not exist, and it may or may not contain various compressed files. To make sure that the directory exists and is empty, you should call the *Static\_Content\_Cache\_Setup* procedure:

```
procedure Static_Content_Cache_Setup
(No_Cache           : in Boolean := False;
 No_Store           : in Boolean := False;
 No_Transform       : in Boolean := False;
 Max_Age            : in AWS.Messages.Delta_Seconds := 86400;
 S_Max_Age          : in AWS.Messages.Delta_Seconds :=
    AWS.Messages.Unset;
 Public             : in Boolean := False;
 Must_Revalidate    : in Boolean := True;
 Proxy_Revalidate   : in Boolean := False);
```

This procedure creates the *Compressed\_Static\_Content\_Cache* directory if it doesn't already exist and if it exists, it deletes everything in it. So basically you're left with an empty directory after a call to this procedure.

The parameters define the content of the *Cache-Control* header sent to the user agent when a request for static content is made. The default parameters are fairly sane, and will result in a cache header looking like this:

```
Cache-Control: max-age=86400, must-revalidate
```

Read the HTTP/1.1 Cache-Control header definition for more information on what the various settings do.

If you use *Yolk.Server* to create and start your AWS server, then this call is made automatically with its default parameters.

If you don't use *Yolk.Server* and you use the *Yolk.Static\_Content* package to serve static content, then you must call this procedure before starting the AWS server. Calling it repeatedly will simply wipe out the compressed cache directory and reset the *Control\_Cache* header according to the given parameters.

## 16.2 `Yolk.Static_Content.Compressable`

If your application contains a bunch of compressable files of significant size, you can save a lot of bandwidth by using the *Compressable* function to serve them:

```
function Compressable
(Request : in AWS.Status.Data)
return AWS.Response.Data;
```

A natural place to call this function would be where you define your content handlers, as seen in this, slightly altered, excerpt from the *Yolk.Handlers* package:

```

declare
  package SC renames Yolk.Static_Content;
  Resource_Handlers : AWS.Services.Dispatchers.URI.Handler;
begin
  AWS.Services.Dispatchers.URI.Register_Regexp
    ( Dispatcher => Resource_Handlers ,
      URI       => "\.css$",
      Action    => Create ( Callback => SC.Compressable 'Access' ));

  AWS.Services.Dispatchers.URI.Register_Regexp
    ( Dispatcher => Resource_Handlers ,
      URI       => "\.html$",
      Action    => Create ( Callback => SC.Compressable 'Access' ));
end;

```

Here you can see that we've defined some content handlers for CSS and HTML files based on a few simple regular expressions, and whenever a resource is requested that matches one of these, the corresponding *Action* callback is called, which in this example is the *Compressable* function. The following steps are then taken:

1. Does the requested resource exist and is it an ordinary file?
  - (a) If no, then return a 404 message
  - (b) if yes, then proceed to 2
2. Does the client support compressed content?
  - (a) If no, then return the requested resource un-compressed
  - (b) if yes, then proceed to 3
3. Is there a compressed version of the resource available on disk?
  - (a) if no, then make one, cache it for future use, and return it
  - (b) If yes, then proceed to 4
4. Is the age of the compressed file  $\leq$  Compressed\_Max\_Age?
  - (a) If no, then delete the current file, build a new one, cache it and return it
  - (b) If yes, then return it

And that's really all there is to it. Note that *Compressable* always looks for requested content in the *WWW\_Root* directory, so if the user requests the file */css/index.css*, then the path is going to be:

```
WWW_Root & "/css/index.ss";
```

Using the default configuration value for *WWW\_Root*, we'll end up with the path:

```
static_content/css/index.css
```

The compressed version of *index.css* is saved as:

```
Compressed_Static_Content_Cache & "/css/index.css" & ".gz";
```

Using the default configuration value for *Compressed\_Static\_Content\_Cache*, we'll end up with the path:

```
static_content/compressed_cache/css/index.css.gz
```

In these two cases the paths are relative to the executable, but you can of course also define *WWW\_Root* and *Compressed\_Static\_Content\_Cache* as absolute paths in the configuration file.

## 16.3 Yolk.Static\_Content.Non\_Compressable

A lot of static content really doesn't benefit from any further compression. This is the case for PNG files, JPEG files and a whole lot of other kinds of files. Static content that has already been compressed, should be handled by the *Non\_Compressable* function:

```
function Non_Compressable  
  (Request : in AWS.Status.Data)  
  return AWS.Response.Data;
```

A call to *Non\_Compressable* is much simpler compared to a call to it's sibling function *Compressable*:

1. Does the requested resource exist and is it an ordinary file?
  - (a) If no, then return a 404 message
  - (b) if yes, then return the file

And that's it. Either the resource is there, or it isn't. If the requested resource is */images/-foo.png*, then the *Non\_Compressable* function searches for the resource in:

```
WWW_Root & "/images/foo.png";
```

Using the default configuration value for *WWW\_Root*, we'll end up with the path:

```
static_content/images/foo.png
```

Just as with *Compressable* a natural place to use *Non\_Compressable* is where you define your content handlers, but it can obviously be used wherever your application handles requests for static content.

## 16.4 Yolk.Static.Set\_Cache\_Options

You can alter the *Cache-Control* header sent to clients requesting static content with a call to *Set\_Cache\_Options*. This is especially handy if you use the *Yolk.Server* package to create and start your server, since the initial call to *Static\_Content\_Cache\_Setup* is hidden you're forced to live with the default values given in the *Static\_Content\_Cache\_Setup* call, resulting in a *Cache-Control* header looking like this:

```
Cache-Control: max-age=86400, must-revalidate
```

If this is not what you want, you can use *Set\_Cache\_Options* to change it:

```
procedure Set_Cache_Options
(No_Cache           : in Boolean := False;
 No_Store           : in Boolean := False;
 No_Transform       : in Boolean := False;
 Max_Age            : in AWS.Messages.Delta_Seconds := 86400;
 S_Max_Age          : in AWS.Messages.Delta_Seconds :=
    AWS.Messages.Unset;
 Public             : in Boolean := False;
 Must_Revalidate    : in Boolean := True;
 Proxy_Revalidate   : in Boolean := False);
```

This **MUST** be called immediately after having called *Yolk.Server.Start*, since the call to *Static\_Content\_Cache\_Setup* in *Start* will overwrite the cache options.

## 17 Yolk.Syndication

The Atom Syndication format (RFC4287) is an XML-based document format that describes lists of related information known as "feeds". Atom documents are used by many web applications as a means of publishing information to users. It's not a complicated format, but it does require a bit of work to construct a proper Atom feed by hand, and since I try my best to avoid work, I made the *Yolk.Syndication* package. This package makes it "easy" to put together an Atom feed document, and have it delivered as a string or an XML/Ada DOM object.

*Yolk.Syndication* helps you construct the Atom XML - it does not check that the resulting XML is valid Atom XML, ie. that the Atom rules have been followed, so if the Atom specification says that

atom:feed elements **MUST** contain exactly one atom:id element.

then it is your job to make sure that your Atom feed has exactly one such element. So before venturing into the realm of Atom, it's probably a good idea to read RFC4287 for a good understanding of the requirements for this document format. A basic example of how to build an Atom feed can be found in the Yolk demo application.

There are two packages in the *Yolk.Syndication* hierarchy:

- *Yolk.Syndication* - Here the necessary types and exceptions are defined.

- *Yolk.Syndication.Writer* - Here the functions and procedures for actually creating an Atom feed are defined.

When time permits, I'll add a *Yolk.Syndication.Reader* package for extraction of data from an Atom feed.

## 17.1 Exceptions

There's one exception in this package:

- *Not\_Valid\_XML*. Is raised when some Xhtml content is not valid XML. This exception can be raised by all procedures that takes Xhtml as content.

## 17.2 The Yolk.Syndication types

There are 5 important types in this package:

- *Text\_Kinds*
- *Relation\_Kinds*
- *Atom\_Entry*
- *Atom\_Entry\_Source*
- *Atom\_Feed*

A few of the procedures in *Yolk.Syndication.Writer* has parameters of the *Text\_Kinds* type. This identifies the kind of data that is being added, with possible values being

- *Text*
- *Html*
- *Xhtml*

Hopefully it's obvious what each of these refers to. Procedures that have parameters of the *Text\_Kinds* type always add data to the feed that can be interpreted as one of these three kinds of text.

The *Relation\_Kinds* type is used in correlation with links that are added to the feed. It identifies how the link relates to the current feed/entry. There are 5 possible relations:

- *Alternate*: Signifies that the link points to an alternate version of the resource described by the containing element.
- *Related*: Signifies that the link points to a resource that is related to, but not the same as, the resource described by the containing element.
- *Self*: Signifies that the link points to a resource that is equivalent to the resource described by the containing element. All feeds should have one link with a *Self* relation pointing to itself.

- *Enclosure*: Signifies that the link points to a resource that is related to, but not the same as, the resource described by the containing element. It also signifies that the resource potentially is large in size and might require special handling. If *Enclosure* is used, then usually the *Length* parameter is set to hint at the size of the resource.
- *Via*: Signifies that the resource provided by the containing element originates in the URI given by the *Href* parameter of the link element.

Finally we have the most important types:

- *Atom\_Entry*: An entry in a feed.
- *Atom\_Entry\_Source*: The origin of an entry.
- *Atom\_Feed*: The Atom feed.

These are core to the functionality of *Yolk.Syndication*, and every single procedure and function in the *Yolk.Syndication.Writer* package use one or the other. These three types are *private*, so the only way to declare an object of either one, is by calling the *New\_Atom\_Feed*, *New\_Atom\_Entry* or *New\_Atom\_Entry\_Source* functions respectively.

Note that only *Atom\_Feed* is thread safe. The two entry related types are not.

## 17.3 Yolk.Syndication.New\_Atom\_Feed

This function initialize an *Atom\_Feed* type, and its specification looks like this:

```
function New_Atom_Feed
  (Base_URI      : in String := None;
   Language      : in String := None;
   Max_Age       : in Duration := 5_616_000.0;
   Max_Entries   : in Positive := 100;
   Min_Entries   : in Positive := 10)
  return Atom_Feed;
```

The *Base\_URI* parameter establish a base for resolving relative references in the feed. The *Language* parameter indicates the natural language used in the feed. *Max\_Age* is a duration that determine when an entry in the feed is old enough to be deleted. *Max\_Entries* is the amount of entries kept in the feed. If there are more than this amount of entries in the feed, then the oldest are deleted until the feed contains *Max\_Entries* entries again. *Min\_Entries* is the minimum amount of entries that must be present in the feed, before we bother deleting entries whose age is > *Max\_Age*. If there are less than *Min\_Entries* entries in the feed, then we keep even a 100 year old entry.

What these parameters hints at, is that some automatic maintenance is done when using *Yolk.Syndication*, and this is indeed true. Usually you do not want a feed to grow forever, and instead of having to manually clear away old stuff, *Yolk.Syndication* handles all this for you, according to the values given when you instantiate an *Atom\_Feed* object using *New\_Atom\_Feed*.



## 17.4 Yolk.Syndication.New\_Atom\_Entry

This function initialize an *Atom\_Entry* type, and its specification looks like this:

```
function New_Atom_Entry
(Base_URI : in String := None;
 Language : in String := None)
return Atom_Entry;
```

As you can see, it's a lot simpler than the *New\_Atom\_Feed* function. This is of course because all the automatic maintenance related parameters have already been set by they *New\_Atom\_Feed* function. All that is left is to define the *Base\_URI* and the natural *Language* used by the entry.

## 17.5 Yolk.Syndication.New\_Atom\_Entry\_Source

This function initialize an *Atom\_Entry\_Source* type, and its specification looks like this:

```
function New_Atom_Entry_Source
(Base_URI : in String := None;
 Language : in String := None)
return Atom_Entry_Source;
```

As you can see, this is very similar to the *New\_Atom\_Entry* function, since its job is to return an *Atom\_Entry\_Source* object, which is basically just an object that describes the origins of a feed entry.

## 17.6 Yolk.Syndication.Writer

In this package we find all the tools necessary to build the Atom XML. There are far too many subprograms in *Yolk.Syndication.Writer* to list here, so instead I'll refer you to the source code. All the subprograms in this package work on one of the *Atom\_Feed*, *Atom\_Entry* or *Atom\_Entry\_Source* types, for example if you want to set a title on your feed, you'll use the *Set\_Title* procedure:

```
procedure Set_Title
(Feed      : in out Atom_Feed;
 Title     : in      String;
 Base_URI  : in      String := None;
 Language  : in      String := None;
 Title_Kind : in      Text_Kinds := Text);
```

Or if the title is for an entry, then:

```
procedure Set_Title
(Entr      : in out Atom_Entry;
 Title     : in      String;
 Base_URI  : in      String := None;
 Language  : in      String := None;
 Title_Kind : in      Text_Kinds := Text);
```

Or we could add an author to the feed/entry/entry source:

```
procedure Add_Author
(Feed      : in out Atom_Feed;
 Name      : in      String;
 Base_URI  : in      String := None;
 Email     : in      String := None;
 Language  : in      String := None;
 URI       : in      String := None);

procedure Add_Author
(Entr      : in out Atom_Entry;
 Name      : in      String;
 Base_URI  : in      String := None;
 Email     : in      String := None;
 Language  : in      String := None;
 URI       : in      String := None);

procedure Add_Author
(Entry_Source : in out Atom_Entry_Source;
 Name         : in      String;
 Base_URI     : in      String := None;
 Email        : in      String := None;
 Language     : in      String := None;
 URI          : in      String := None);
```

There's actually a hint about the kind of XML element that is being produced by these procedures. If the name of the procedure starts with *Set\_* then it hints at an XML element of which there's only ever one. So *Set\_Title* creates a `<title>Foo</title>` element, and if called again, overwrites the previous value, whereas if the name of procedure starts with *Add\_* then the Atom specification allows for multiples of these, as can be seen with the *Add\_Author* procedure. Calling this one creates an `<author>` element, and calling it again simply adds one more author to the feed/entry/entry source.

Most of the subprograms in this package deals directly with building the Atom XML, but there are a few exceptions, as you will see next.

## 17.7 Counting the amount of entries in a feed

If you need to know how many entries that are currently in a feed, the you can use the *Yolk.Syndication.Writer.Amount\_Of\_Entries* function:

```
function Amount_Of_Entries
(Feed : in Atom_Feed)
return Natural;
```

## 17.8 Clearing and deleting entries

There are two procedures available for clearing and deleting entries, one for clearing all entries away, and one for deleting entries based on their Id. Lets start with the one that clears out everything:

```
procedure Clear_Entry_List
  (Feed : in out Atom_Feed);
```

Calling *Clear\_Entry\_List* deletes every single entry that has been added to the *Feed* object so far. A less destructive procedure is *Delete\_Entry*:

```
procedure Delete_Entry
  (Feed : in out Atom_Feed;
   Id    : in      String);
```

Using this one, you can delete all entries whose *Id* is *Id*. Note that the match must be exact, so case matters. FOO is not the same as foo.

## 17.9 Getting the Atom feed

When you've added titles, authors, categories, entries and content to your feed, the next step is turning it into XML. This can be done in one of two ways:

1. As string XML.
2. As an XML/Ada DOM XML object.

The string XML is obviously for the final audience, whereas the DOM XML is useful if you need to do some further work on the feed before releasing it on the world. Here they are:

```
function Get_XML_DOM
  (Feed : in Atom_Feed)
  return DOM.Core.Document;

function Get_XML_String
  (Feed           : in Atom_Feed;
   Pretty_Print  : in Boolean := False)
  return String;
```

Note that if *Pretty\_Print* is *True* then whitespace is going get mangled according to these rules:

If *Pretty\_Print* is true, then the XML nodes will be indented so that children nodes are to the right of their parents. It is set to False by default because its use changes the document (addition or removal of whitespaces among other things), which in general has no effect for automatic tools reading the document. All whitespaces are modified outside of elements containing nothing but text nodes. For text nodes, leading and trailing whitespaces are also deleted.

*Get\_XML\_String* relies on the *Write* function from XML/Ada to generate its output, and the above quote is taken straight from the XML/Ada source comments. It's also worth noting that *Get\_XML\_String* and *Get\_XML\_DOM* both are pretty resource-hungry, so it's probably best to cache the results for later use, instead of calling them on each and every hit.

## 18 Yolk.Utilities

This package contains various support functionality.

## 19 Yolk.Whoops

This package contains one single procedure: *Unexpected\_Exception\_Handler*. I suspect the name gives away exactly what this procedure does. It looks like this:

```
procedure Unexpected_Exception_Handler
(E      : Ada.Exceptions.Exception_Occurrence ;
Log     : in out AWS.Log.Object ;
Error   : AWS.Exceptions.Data ;
Answer  : in out AWS.Response.Data ) ;
```

You can use this procedure to catch any and all exceptions you've failed to catch in your application, ie. the ones AWS pickup and fail to do anything sensible with. When *Unexpected\_Reception\_Handler* catch an exception, two things happen:

1. The exception is logged to the *Yolk.Log.Error* trace.
2. A HTTP status code 500 is returned to the user.

The template used to create the HTTP code 500 error message is *demo/exe/templates/system/500.tmpl*. You can of course change this to match the look and feel of your application. The path to the *500.tmpl* file is set by the configuration setting *System\_Templates\_Path*.

You can see an example on how to use this procedure in the file *demo/src/yolk\_demo.adb*.